

Table of Contents

1. Introduction
 2. Phase 1: Data Ingestion
 3. Phase 2: Data Cleaning
 4. Phase 3: Data Validation
 5. Phase 4: Network KPIs
 6. Phase 5: Customer Experience
 7. Phase 6: Window Functions
 8. Phase 7: Anomaly Detection
-

CASE STUDY 7

Title: Telecom Network Quality and Customer Experience Analytics using PySpark

Spark Session Initialization

```
# Initialize Spark Session
from pyspark.sql import SparkSession

spark = SparkSession.builder \
    .appName("TelecomNetworkQuality") \
    .getOrCreate()
```

Required Imports

```
from pyspark.sql.functions import (
    col, trim, upper, lower, when, regexp_replace,
    avg, count, sum as _sum,
    try_to_timestamp, coalesce,
    row_number, lag,
    round, expr
)
from pyspark.sql.window import Window
from pyspark.sql.types import StructType, StructField, StringType, IntegerType,
    DoubleType
```

PHASE 1: INGESTION

1. Define Schema and Read CSV

```
# 1. Read network_logs.csv as all StringType
schema_string = StructType([
    StructField("event_id", StringType(), True),
    StructField("subscriber_id", StringType(), True),
    StructField("tower_id", StringType(), True),
    StructField("city", StringType(), True),
    StructField("network_type", StringType(), True),
    StructField("signal_strength", StringType(), True),
    StructField("download_speed_mbps", StringType(), True),
    StructField("upload_speed_mbps", StringType(), True),
    StructField("latency_ms", StringType(), True),
    StructField("call_drop", StringType(), True),
    StructField("event_time", StringType(), True),
    StructField("device_type", StringType(), True)
])

df_raw = spark.read.csv("/content/network_logs.csv", header=True,
schema=schema_string)
```

2. Print Schema and Row Count

```
# 2. Print schema and row count
df_raw.printSchema()
print(f"Raw Row Count: {df_raw.count()}")
```

```
root
|-- event_id: string (nullable = true)
|-- subscriber_id: string (nullable = true)
|-- tower_id: string (nullable = true)
|-- city: string (nullable = true)
|-- network_type: string (nullable = true)
|-- signal_strength: string (nullable = true)
|-- download_speed_mbps: string (nullable = true)
|-- upload_speed_mbps: string (nullable = true)
|-- latency_ms: string (nullable = true)
|-- call_drop: string (nullable = true)
|-- event_time: string (nullable = true)
|-- device_type: string (nullable = true)
```

Raw Row Count: 180000

3. Show Sample Rows

```
# 3. Show sample rows
df_raw.show(5)
```

```
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
|event_id|subscriber_id|tower_id|
city|network_type|signal_strength|download_speed_mbps|upload_speed_mbps|latency_ms|call_dr
op|          event_time| device_type|
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
| E100000|      S5975|     T837|Bangalore|       3G|      invalid|
NULL|        NULL| invalid|      YES|01/01/2026 00:00:00|    Android|
| E100001|      S3537|     T283|Hyderabad|       5G|      -83|
124.07|        41.26|      114|      NO|2026-01-01 00:00:03|FeaturePhone|
| E100002|      S1629|     T877|      Pune|       4G|      -72|
41.01|        3.36|      221|      NO|2026-01-01 00:00:06|FeaturePhone|
| E100003|      S9422|     T431|     Delhi|       3G|      -97|
46.98|        13.36|      148|      NO|2026-01-01 00:00:09|    Android|
| E100004|      S1776|     T432|Hyderabad|       3G|      -83|
15.3|        31.1|      251|      NO|2026-01-01 00:00:12|FeaturePhone|
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
```

PHASE 2: CLEANING

1. Trim String Columns

```
# 1. Trim string columns
df_trimmed = df_raw.select([trim(col(c)).alias(c) for c in df_raw.columns])

df_trimmed.show(5)
```

```
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
|event_id|subscriber_id|tower_id|
city|network_type|signal_strength|download_speed_mbps|upload_speed_mbps|latency_ms|call_dr
op|          event_time| device_type|
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+
| E100000|      S5975|     T837|Bangalore|       3G|      invalid|
NULL|        NULL| invalid|      YES|01/01/2026 00:00:00|      Android|
| E100001|      S3537|     T283|Hyderabad|       5G|        -83|
124.07|        41.26|      114|        NO|2026-01-01 00:00:03|FeaturePhone|
| E100002|      S1629|     T877|      Pune|       4G|        -72|
41.01|        3.36|      221|        NO|2026-01-01 00:00:06|FeaturePhone|
| E100003|      S9422|     T431|      Delhi|       3G|        -97|
46.98|        13.36|      148|        NO|2026-01-01 00:00:09|      Android|
| E100004|      S1776|     T432|Hyderabad|       3G|        -83|
15.3|        31.1|      251|        NO|2026-01-01 00:00:12|FeaturePhone|
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+
```

2. Normalize String Fields

```
# 2. Normalize string fields (city, network_type, device_type, call_drop)
df_normalized = df_trimmed.withColumn("city", upper(col("city")))
    .withColumn("network_type", upper(col("network_type")))
    .withColumn("device_type", upper(col("device_type")))
    .withColumn("call_drop", upper(col("call_drop")))

df_normalized.show(5)
```

| event_id | subscriber_id | tower_id | city | network_type | signal_strength | download_speed_mbps | upload_speed_mbps | latency_ms | call_drop | event_time | device_type |
|----------|---------------|----------------|------|---------------------|-----------------|---------------------|-------------------|------------|--------------|------------|-------------|
| E100000 | S5975 | T837 BANGALORE | | 3G | | | | | invalid | | |
| NULL | NULL | invalid | YES | 01/01/2026 00:00:00 | | | | | ANDROID | | |
| E100001 | S3537 | T283 HYDERABAD | | 5G | | | | | -83 | | |
| 124.07 | 41.26 | 114 | NO | 2026-01-01 00:00:03 | | | | | FEATUREPHONE | | |
| E100002 | S1629 | T877 PUNE | | 4G | | | | | -72 | | |
| 41.01 | 3.36 | 221 | NO | 2026-01-01 00:00:06 | | | | | FEATUREPHONE | | |
| E100003 | S9422 | T431 DELHI | | 3G | | | | | -97 | | |
| 46.98 | 13.36 | 148 | NO | 2026-01-01 00:00:09 | | | | | ANDROID | | |
| E100004 | S1776 | T432 HYDERABAD | | 3G | | | | | -83 | | |
| 15.3 | 31.1 | 251 | NO | 2026-01-01 00:00:12 | | | | | FEATUREPHONE | | |

3. Clean Numeric Fields Safely

```

# 3. Clean numeric fields safely
df_cleaned = (
    df_normalized
    .withColumn(
        "signal_strength",
        when(
            regexp_replace(col("signal_strength"), "[^0-9-]", "") == "",
            None
        ).otherwise(
            regexp_replace(col("signal_strength"), "[^0-9-]", ""
        ))
    )
    .withColumn(
        "download_speed_mbps",
        when(
            regexp_replace(col("download_speed_mbps"), "[^0-9.]", "") == "",
            None
        ).otherwise(
            regexp_replace(col("download_speed_mbps"), "[^0-9.]",
        ))
    )
    .withColumn(
        "upload_speed_mbps",
        when(
            regexp_replace(col("upload_speed_mbps"), "[^0-9.]", "") == "",
            None
        ).otherwise(
            regexp_replace(col("upload_speed_mbps"), "[^0-9.]",
        ))
    )
    .withColumn(
        "latency_ms",
        when(
            regexp_replace(col("latency_ms"), "[^0-9]", "") == "",
            None
        ).otherwise(
            regexp_replace(col("latency_ms"), "[^0-9]", "").cast(IntegerType())
        )
    )
)

df_cleaned.show(5)

```

```

+-----+-----+-----+-----+-----+
|event_id|subscriber_id|tower_id|
city|network_type|signal_strength|download_speed_mbps|upload_speed_mbps|latency_ms|call_dr
op|          event_time| device_type|
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+
| E100000|      S5975|    T837|BANGALORE|       3G|        NULL|
NULL|      NULL|      NULL|      YES|01/01/2026 00:00:00|      ANDROID|
| E100001|      S3537|    T283|HYDERABAD|       5G|       -83|
124.07|      41.26|     114|       NO|2026-01-01 00:00:03|FEATUREPHONE|
| E100002|      S1629|    T877|      PUNE|       4G|       -72|
41.01|      3.36|     221|       NO|2026-01-01 00:00:06|FEATUREPHONE|
| E100003|      S9422|    T431|      DELHI|       3G|       -97|
46.98|      13.36|     148|       NO|2026-01-01 00:00:09|      ANDROID|
| E100004|      S1776|    T432|HYDERABAD|       3G|       -83|
15.3|      31.1|     251|       NO|2026-01-01 00:00:12|FEATUREPHONE|
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+

```

4. Parse Event Time with Multiple Formats

```

# 4. Parse event_time with multiple formats
df = df_cleaned.withColumn(
    "event_time_clean",
    expr("""
        coalesce(
            try_to_timestamp(event_time, 'yyyy-MM-dd HH:mm:ss'),
            try_to_timestamp(event_time, 'dd/MM/yyyy HH:mm:ss'),
            try_to_timestamp(event_time, 'yyyy/MM/dd HH:mm:ss')
        )
    """)
)

df.show(5)

```

```

+-----+-----+-----+-----+-----+
|event_id|subscriber_id|tower_id|
+-----+-----+-----+-----+
|city|network_type|signal_strength|download_speed_mbps|upload_speed_mbps|latency_ms|call_drop|event_time|device_type|event_time_clean|
+-----+-----+-----+-----+-----+
| E100000|      S5975|     T837|BANGALORE|        3G|       NULL|
NULL|      NULL|      NULL|      YES|01/01/2026 00:00:00|      ANDROID|2026-01-01
00:00:00|
| E100001|      S3537|     T283|HYDERABAD|        5G|       -83|
124.07|      41.26|      114|      NO|2026-01-01 00:00:03|FEATUREPHONE|2026-01-01
00:00:03|
| E100002|      S1629|     T877|      PUNE|        4G|       -72|
41.01|      3.36|      221|      NO|2026-01-01 00:00:06|FEATUREPHONE|2026-01-01
00:00:06|
| E100003|      S9422|     T431|      DELHI|        3G|       -97|
46.98|      13.36|      148|      NO|2026-01-01 00:00:09|      ANDROID|2026-01-01
00:00:09|
| E100004|      S1776|     T432|HYDERABAD|        3G|       -83|
15.3|      31.1|      251|      NO|2026-01-01 00:00:12|FEATUREPHONE|2026-01-01
00:00:12|
+-----+-----+-----+-----+-----+
|-----+-----+-----+-----+-----+

```

PHASE 3: VALIDATION

1. Count Invalid Values for Each Numeric Field

```
# 1. Count invalid values for each numeric field
invalid_counts = df.select(
    _sum(col("signal_strength").isNull().cast("int")).alias("invalid_signal_strength"),
    _sum(col("download_speed_mbps").isNull().cast("int")).alias("invalid_download_speed"),
    _sum(col("upload_speed_mbps").isNull().cast("int")).alias("invalid_upload_speed"),
    _sum(col("latency_ms").isNull().cast("int")).alias("invalid_latency")
)
print("Invalid Value Counts:")
invalid_counts.show()
```

Invalid Value Counts:

| invalid_signal_strength | invalid_download_speed | invalid_upload_speed | invalid_latency |
|-------------------------|------------------------|----------------------|-----------------|
| 9474 | 13764 | 5807 | 4865 |

2. Count Invalid Timestamps

```
# 2. Count invalid timestamps
print("Invalid timestamps:")
df.filter(col("event_time_clean").isNull()).count()
```

Invalid timestamps:

1605

3. Remove Duplicate Logs

```
# 3. Remove duplicate logs  
df = df.dropDuplicates(["event_id"])  
print(f"Count after removal: {df.count()}")
```

Count after removal: 180000

PHASE 4: NETWORK KPIs

1. Average Download Speed per City

```
# 1. Average download speed per city
avg_download_city = df.groupBy("city") \
    .agg(round(avg("download_speed_mbps"), 2).alias("avg_download_speed"))

avg_download_city.show()
```

| city | avg_download_speed |
|-----------|--------------------|
| KOLKATA | 75.83 |
| DELHI | 75.79 |
| BANGALORE | 75.49 |
| HYDERABAD | 75.4 |
| CHENNAI | 75.51 |
| PUNE | 75.33 |
| MUMBAI | 75.53 |

2. Average Latency per City

```
# 2. Average latency per city
avg_latency_city = df.groupBy("city") \
    .agg(round(avg("latency_ms"), 2).alias("avg_latency"))

avg_latency_city.show()
```

```
+-----+-----+
|   city|avg_latency|
+-----+-----+
| KOLKATA|      154.95|
| DELHI |      155.44|
|BANGALORE|      156.1|
|HYDERABAD|      155.13|
| CHENNAI|      154.65|
|     PUNE|      154.94|
| MUMBAI|      154.44|
+-----+-----+
```

3. Call Drop Rate per City

```
# 3. Call drop rate per city
call_drop_city = df.groupBy("city") \
    .agg(
        round(
            _sum(when(col("call_drop") == "YES", 1).otherwise(0)) / count("*"),
            3
        ).alias("call_drop_rate")
    )

call_drop_city.show()
```

```
+-----+-----+
|   city|call_drop_rate|
+-----+-----+
| KOLKATA|      0.05|
| DELHI |      0.051|
|BANGALORE|      0.05|
|HYDERABAD|      0.051|
| CHENNAI|      0.049|
|     PUNE|      0.049|
| MUMBAI|      0.051|
+-----+-----+
```

4. Call Drop Rate per Tower

```

# 4. Call drop rate per tower
call_drop_tower = df.groupBy("tower_id") \
    .agg(
        round(
            _sum(when(col("call_drop") == "YES", 1).otherwise(0)) / count("*"),
            3
        ).alias("call_drop_rate")
    )
call_drop_tower.show(5)

```

| tower_id | call_drop_rate |
|----------|----------------|
| T352 | 0.051 |
| T929 | 0.027 |
| T947 | 0.038 |
| T590 | 0.063 |
| T855 | 0.057 |

only showing top 5 rows

5. Top 10 Worst Towers

```

# 5. Top 10 worst towers
# Criteria: High Drop Rate, High Latency, Low Download Speed
worst_towers = df.groupBy("tower_id") \
    .agg(
        round(avg("latency_ms"), 2).alias("avg_latency"),
        round(avg("download_speed_mbps"), 2).alias("avg_download_speed"),
        round(
            _sum(when(col("call_drop") == "YES", 1).otherwise(0)) / count("*"),
            3
        ).alias("call_drop_rate")
    ) \
    .orderBy(col("call_drop_rate").desc(),
            col("avg_latency").desc(),
            col("avg_download_speed").asc()) \
    .limit(10)

worst_towers.show()

```

| tower_id | avg_latency | avg_download_speed | call_drop_rate |
|----------|-------------|--------------------|----------------|
| T358 | 162.98 | 75.27 | 0.104 |
| T275 | 149.97 | 75.14 | 0.1 |
| T241 | 152.49 | 72.87 | 0.097 |
| T455 | 156.39 | 77.05 | 0.096 |
| T697 | 150.16 | 80.92 | 0.092 |
| T538 | 160.16 | 72.71 | 0.089 |
| T257 | 155.76 | 79.97 | 0.089 |
| T653 | 151.88 | 78.33 | 0.088 |
| T157 | 155.75 | 77.22 | 0.087 |
| T659 | 153.3 | 81.18 | 0.087 |

PHASE 5: CUSTOMER EXPERIENCE

Compute Metrics for Each Subscriber

```
# Compute metrics for each subscriber_id
subscriber_metrics = df.groupBy("subscriber_id") \
    .agg(
        count("*").alias("event_count"),
        round(avg("download_speed_mbps"), 2).alias("avg_download_speed"),
        round(avg("latency_ms"), 2).alias("avg_latency"),
        _sum(when(col("call_drop") == "YES",
1).otherwise(0)).alias("call_drop_count")
    )

subscriber_metrics.show(5)
```

| subscriber_id | event_count | avg_download_speed | avg_latency | call_drop_count |
|---------------|-------------|--------------------|-------------|-----------------|
| S2422 | 22 | 76.5 | 156.36 | 0 |
| S1828 | 26 | 61.47 | 120.62 | 0 |
| S2414 | 14 | 73.25 | 142.62 | 0 |
| S7616 | 17 | 95.41 | 158.35 | 0 |
| S6467 | 18 | 80.4 | 174.33 | 1 |

only showing top 5 rows

Poor Experience Users

```
# Poor experience logic
# Logic defined: High drops (>=3), Low speed (<5 Mbps), High latency (>200ms)
poor_experience_users = subscriber_metrics.filter(
    (col("call_drop_count") >= 3) |
    (col("avg_download_speed") < 5) |
    (col("avg_latency") > 200)
)

poor_experience_users.show(5)
```

| subscriber_id | event_count | avg_download_speed | avg_latency | call_drop_count |
|---------------|-------------|--------------------|-------------|-----------------|
| S7972 | 21 | 74.89 | 150.95 | 5 |
| S7123 | 20 | 59.2 | 145.6 | 3 |
| S2497 | 18 | 92.98 | 138.06 | 3 |
| S4911 | 27 | 71.12 | 121.85 | 3 |
| S2181 | 23 | 77.06 | 132.18 | 3 |

only showing top 5 rows

PHASE 6: WINDOW FUNCTIONS

1. Rank Towers Within Each City by Call Drop Rate

```
# 1. Rank towers within each city by call drop rate
tower_city_window =
Window.partitionBy("city").orderBy(col("call_drop_rate").desc())

tower_ranked = call_drop_city.withColumn(
    "tower_rank",
    row_number().over(tower_city_window)
)
tower_ranked.show(5)
```

```
+-----+-----+-----+
|     city|call_drop_rate|tower_rank|
+-----+-----+-----+
|BANGALORE|      0.05|      1|
| CHENNAI|      0.049|      1|
|   DELHI|      0.051|      1|
|HYDERABAD|      0.051|      1|
| KOLKATA|      0.05|      1|
+-----+-----+-----+
only showing top 5 rows
```

2. Rank Subscribers Within Each City by Worst Experience

```

# 2. Rank subscribers within each city by worst experience
subscriber_city_df = df.join(subscriber_metrics, "subscriber_id")

subscriber_window = Window.partitionBy("city").orderBy(
    col("call_drop_count").desc(),
    col("avg_latency").desc(),
    col("avg_download_speed").asc()
)

subscriber_ranked = subscriber_city_df.withColumn(
    "experience_rank",
    row_number().over(subscriber_window)
)
subscriber_ranked.show()

```

| subscriber_id | event_id | tower_id | city | network_type | signal_strength | download_speed_mbps | upload_speed_mbps | latency_ms | call_drop |
|---------------|----------------|----------------|------|--------------|-----------------|---------------------|-------------------|------------|-----------|
| 78.83 | S8721 E100973 | T662 BANGALORE | | 3G | 5.34 | 101 | NO | -77 | |
| 86.53 | S8721 E105340 | T581 BANGALORE | | 5G | 48.66 | 226 | YES | -89 | |
| 116.09 | S8721 E108460 | T131 BANGALORE | | 4G | 29.66 | 286 | YES | -82 | |

[Output truncated for brevity]

3. Signal Deterioration Detection Using Lag

```

# 3. Signal deterioration detection using lag
signal_window = Window.partitionBy("tower_id").orderBy("event_time_clean")

df_signal_lag = df.withColumn(
    "prev_signal_strength",
    lag("signal_strength").over(signal_window)
).withColumn(
    "signal_drop",
    col("prev_signal_strength") - col("signal_strength")
)

df_signal_lag.show()

```

| event_id | subscriber_id | tower_id | city | network_type | signal_strength | download_speed_mbps | upload_speed_mbps | latency_ms | call_drop |
|----------|---------------|----------|-----------|--------------|-----------------|---------------------|-------------------|------------|-----------|
| E100028 | S9371 | T102 | BANGALORE | 5G | 52.78 | 45.88 | 285 | NO | -62 |
| E101314 | S5759 | T102 | MUMBAI | 3G | 38.74 | 17.97 | 278 | NO | -81 |

[Output truncated for brevity]

PHASE 7: ANOMALY DETECTION

Detect Towers with Anomalies

```
# Detect towers where: Latency spikes, Download speed drops, Call drops spikes
rolling_window = Window.partitionBy("tower_id") \
    .orderBy("event_time_clean") \
    .rowsBetween(-3, -1)

df_anomaly = df.withColumn(
    "rolling_avg_latency",
    avg("latency_ms").over(rolling_window)
).withColumn(
    "rolling_avg_download",
    avg("download_speed_mbps").over(rolling_window)
).withColumn(
    "latency_spike",
    col("latency_ms") > col("rolling_avg_latency") * 1.5
).withColumn(
    "download_drop",
    col("download_speed_mbps") < col("rolling_avg_download") * 0.5
).withColumn(
    "call_drop_spike",
    col("call_drop") == "YES"
)

anomalous_events = df_anomaly.filter(
    col("latency_spike") | col("download_drop") | col("call_drop_spike")
)

anomalous_events.show()
```

```

+-----+-----+-----+-----+-----+
|event_id|subscriber_id|tower_id|
city|network_type|signal_strength|download_speed_mbps|upload_speed_mbps|latency_ms|call_dr
op|
+-----+-----+-----+-----+-----+
| E103037|      S7132|    T102|HYDERABAD|        5G|       -92|
27.06|          4.23|        257|        NO|
| E103866|      S2572|    T102| KOLKATA|        3G|       -81|
31.06|          44.33|        168|        NO|
| E104673|      S3570|    T102|HYDERABAD|        5G|       -72|
18.78|          42.25|        118|        NO|

```

[Output truncated for brevity]

End of Document

Telecom Network Quality and Customer Experience Analytics using PySpark